

# Lab 3: BST Comparison:

## Performance analysis and detailed approach

Qiyong Wu

October 28, 2024

## Abstract

In this lab, we explore the concurrent computation of binary search tree (BST) hashes and tree comparison using Go's concurrency tools, specifically goroutines, channels, and synchronization primitives. The goal is to optimize the processes of hashing, identifying duplicate hash groups, and performing pairwise tree comparisons to achieve high efficiency while maintaining accuracy. Each BST is represented by a unique hash generated from an in-order traversal using a custom hashing function. Multiple configurations of worker threads handle hash computation and map updates, providing insights into the impact of parallelism on performance. This report evaluates the scalability of different parallelization approaches across several worker configurations, with a focus on balancing computational and synchronization overhead.

## Introduction

Binary Search Trees (BSTs) are a fundamental data structure in computer science, used extensively in applications requiring efficient data retrieval and organization. This lab focuses on implementing a parallelized approach to compare BSTs using a custom hash-based methodology. Given a set of BSTs, we aim to compute a unique hash for each tree using an in-order traversal, group trees with identical hashes, and further compare these groups to identify structurally identical trees. Leveraging Go's concurrency model, the implementation employs goroutines and channels for inter-process communication and synchronization mechanisms such as mutexes and semaphores to ensure data consistency during concurrent hash computations and updates.

The lab is structured in three parts:

- **Hash Computation:** We compute the hash of each BST in parallel by spawning a designated number of worker goroutines, based on the `-hash-workers` flag. This step evaluates the efficiency of parallel hash computation and measures the time taken to generate hashes for all BSTs.

- **Hash Grouping:** Trees with identical hashes are grouped to identify potential duplicates. Different configurations of data worker goroutines are managed by the `-data-workers` flag to control access to the map holding hash groups. This part evaluates synchronization overhead and seeks to balance hash computation with data aggregation.
- **Tree Comparison:** For each group with duplicate hashes, tree comparisons are performed in parallel using `-comp-workers`. This step refines the grouping by identifying structurally identical trees, based on a pairwise comparison strategy.

The experimental setup uses multiple input files with varying BST complexities to analyze the performance of each concurrent approach. The results are assessed for both correctness and performance under different configurations, providing insights into how concurrency levels and synchronization impact the speed and scalability of BST hashing and comparison.

This report presents the design and implementation of these different approaches, along with performance analyses based on their execution times for various datasets. We also discuss the challenges encountered during the optimization process, including memory management and non-deterministic behavior caused by atomic operations in CUDA. By comparing the performance of these implementations, we aim to highlight the trade-offs between different levels of abstraction in GPU programming and the impact of hardware-level optimizations on the performance of parallel algorithms like KMeans.

# 1 Background and Hashing Methodology

## 1.1 Binary Search Tree Structure

My implementation of binary search tree (BST) structure is defined by two key components:

- **TreeNode:** Represents a node in the BST with the following fields:
  - `Val` - The integer value stored in the node.
  - `Left` - Pointer to the left child (subtree).
  - `Right` - Pointer to the right child (subtree).
  - `Parent` - (Optional) Pointer to the parent node.
- **BinarySearchTree:** This structure has the following fields:
  - `Id` - A unique identifier for the BST.
  - `Root` - Pointer to the root `TreeNode` of the BST.
  - `Hash` - A hash value representing the unique structure and content of the tree.
  - `InOrderTraversal` - A list storing node values in in-order sequence for comparison and hashing.

Operations including:

- **Insertion:** New values are inserted into the tree following BST rules: values smaller than the current node go to the left child, and values larger go to the right child.
- **Traversal:** The in-order traversal function captures all node values in sorted order. This sequence is used for generating a hash, which enables efficient comparison and identification of equivalent trees.

## 1.2 Hash Computation Method

1. **In-Order Traversal:** If the `InOrderTraversal` field of the BST is empty, the function `ComputeHash` populates it by performing an in-order traversal of the tree. This traversal arranges the node values in sorted order.
2. **Hash Calculation:** The hash value is initialized to 1. For each value  $v$  in the in-order traversal, the following steps are applied:
  - Compute a transformed value,  $\text{newValue} = v + 2$ .
  - Update the hash using the formula:

$$\text{hash} = (\text{hash} \times \text{newValue} + \text{newValue}) \mod 1000$$

This transformation and modulus operation ensures the hash remains within a bounded range (0 to 999).

3. **Output:** The computed hash uniquely represents the sequence of values in the BSTs in-order traversal, enabling efficient equivalence checks between trees. Trees with different structures but identical contents yield the same hash.

## 2 Implementation

### 2.1 Hash Computation

We are using the hash function provided as above to calculate the hash of a BST and compare them according to ids in same hash group.

#### 2.1.1 Sequential Hash Computation

In the sequential implementation:

- **Single-Threaded Execution:** Each binary search tree (BST) is processed sequentially in a single thread.
- **Direct Hash Mapping:** For each BST, a hash is computed and stored immediately in a map (`hash2bstId`) without concurrency handling.
- **Performance:** This approach has minimal overhead but lacks parallelism, limiting speed on large datasets.

### 2.1.2 Parallel Hash Computation

Parallel hash computation is implemented in three variants:

#### 1. Channel-Based Implementation:

- **Hash Workers** compute hashes concurrently and send results (BST ID and hash) through a channel.
- **Data Workers** consume from the channel and update the shared map with a mutex to ensure thread safety.

#### 2. Single-Mutex Implementation:

- Each hash worker both computes the hash and directly updates the shared map.
- A single mutex guards access to the map, simplifying concurrency but risking contention under high load.

#### 3. Semaphore-based implementation with Fine-Grained Locks:

- The map is partitioned into shards, each with its own lock.
- Hash workers compute hashes and update only relevant shards, reducing contention and improving scalability on larger datasets.

## 2.2 Hash Grouping

### 2.2.1 UnionFind data structure for grouping

Components of Union-Find:

- **parent []int:** An array where **parent[i]** represents the parent ID of BST *i*. Initially, each BST is its own parent, forming individual groups.
- **rank []int:** Tracks the rank (or depth) of each subtree to balance the tree during union operations, favoring higher rank trees as roots.
- **locks map[int]\*sync.Mutex:** A map of mutex locks for each BST ID to ensure thread-safe access during concurrent union operations.
- **mu sync.Mutex:** A global mutex that manages access to the **locks** map, allowing safe, on-demand lock creation.

Key Functions:

- **NewUnionFind(size int):** Initializes the **UnionFind** structure, setting each BST as its own parent and assigning an empty rank.
- **getLock(id int):** Returns a lock for a given ID, creating it if it does not exist, with **mu** ensuring thread-safe additions to the **locks** map.

- **Find(x int):** Implements *path compression* by updating the parent of each node in the path to point directly to the root, reducing future access times by flattening the tree structure.
- **Union(x, y int):** Joins two BST groups represented by IDs x and y.
  - Locks IDs in a consistent order to avoid deadlocks.
  - Finds the root parents of both IDs and merges them based on **rank**:
    - \* If **rank[rootX] < rank[rootY]**, **rootX** is attached under **rootY**.
    - \* If **rank[rootX] > rank[rootY]**, **rootY** is attached under **rootX**.
    - \* If ranks are equal, **rootY** is attached under **rootX**, and **rank[rootX]** is incremented.

The **UnionFind** structure groups BSTs with identical hashes (equivalent trees) by linking their IDs. During comparison, the **Union** operation connects equivalent BST IDs, and after all comparisons, **Find** retrieves the root ID for each group. This design enables efficient and concurrent grouping of equivalent BSTs, essential for handling large datasets with parallel comparisons.

### 2.2.2 Sequential Hash Grouping (**channelImpl**)

In the sequential approach, hash grouping is handled by the **sequentialImpl** function:

- **Hash Computation:** Each binary search tree (BST) in **bstList** is processed sequentially in a single thread.
- **Direct Mapping:** For each BST, a hash is computed using **ComputeHash**, and the hash is used as a key in a map (**hash2bstId**). Each key points to a list of BST IDs (indices) that share the hash.
- **No Concurrency Handling:** Since the map is accessed in a single-threaded manner, no additional concurrency controls are needed.

### 2.2.3 Parallel Hash Grouping

The parallel implementations use different methods to achieve hash grouping. First we need to initialize a map (**hash2bstId**) to store tree indices by hash value, facilitating efficient grouping and then the remaining hashing is done by different implementations depends on the data-workers and hash-workers value.

#### Channel-Based Implementation (**channelImpl**)

- **Worker Channels:** A channel (**taskCh**) assigns tasks to multiple **hashWorkers**, each responsible for computing a hash for a different BST. Each worker sends its result (BST ID and hash) to a result channel (**resultCh**).

- **Data Workers:** Multiple `dataWorkers` consume from `resultCh`. Each worker uses a mutex to safely update the shared map `hash2bstId`. A lock (`hashLocksMu`) ensures that each hash has a unique mutex in `hashLocks`, used to guard entries when appending BST IDs.

### Single-Mutex Implementation (`mutexImpl`)

- **Mutex for Synchronization:**
  - Uses a single `sync.Mutex` (`mu`) to ensure that only one goroutine can update the map at a time, thus preventing data races.
- **Concurrent Hash Computation:**
  - Distributes tasks among multiple `hashWorkers` through a channel. Each worker computes hashes and updates the shared map under mutex protection.
- **Performance Considerations:**
  - The use of a single mutex simplifies the concurrency model but may lead to contention, limiting scalability when the number of workers is high.

### Semaphore-based Implementation (`semaphoreImpl`)

- **Sharded Map:**
  - Divides the map into several segments, each with its own lock, allowing more granular concurrency and reducing contention significantly.
- **Efficient Shard Management:**
  - Assigns hashes to shards using a modulus operation, ensuring a balanced distribution of data across shards.
- **Result Merging:**
  - Combines data from all shards into a single map at the end of processing, preparing it for the next steps of tree comparison.

## 2.3 Tree Comparison

### 2.3.1 Sequential Tree Comparison

In the sequential approach to tree comparison:

- **Hash Check:** Initially, each BST pair within the same hash group is compared by their hash values. If the hashes are different, the trees are immediately deemed non-equivalent.

- **Direct Comparison:** For BST pairs that share the same hash, a detailed comparison is performed using the `CompareBST` function. This function checks the in-order traversal of the trees for equivalence. The results are recorded in an adjacency matrix, indicating equivalence between trees.

### 2.3.2 Parallel Tree Comparison

In the parallel approach to tree comparison:

- **Comparison Channel:** Pairs of BST IDs from hash groups with identical hashes are sent through a channel (`compCh`) to worker goroutines dedicated to comparisons.
- **Worker Pool:** A pre-determined number of comparison workers (`compWorkers`) retrieve pairs from `compCh`, conduct the comparisons using `CompareBST`, and update an adjacency matrix based on whether the trees are structurally identical.
- **Synchronization:** A `sync.WaitGroup` is employed to ensure all comparisons are completed before the process moves forward, maintaining data integrity and preventing concurrent access issues.

## Experimental Setup

### 2.4 Input Data and Flags

The BST equivalence program uses the following flags to control input and synchronization:

- `-input=<path>`: Specifies the path to the input file with binary search trees (BSTs), such as `simple.txt`, `coarse.txt`, and `fine.txt`.
- `-hash-workers=<number>`: Sets the number of goroutines for hashing the BSTs, enabling parallel computation across threads.
- `-data-workers=<number>`: Determines the number of goroutines updating the hash map. The combination of `hash-workers` and `data-workers` defines the synchronization approach.

### 2.5 Synchronization Strategies

The program adapts its synchronization mechanism based on the following flag combinations:

- `-hash-workers=1 -data-workers=1`: Sequential mode, where hashing and map updates occur in the main thread without concurrency.
- `-hash-workers=i -data-workers=1` (**where**  $i > 1$ ):  $i$  goroutines compute hashes and send results to a central manager via a channel for map updates.
- `-hash-workers=i -data-workers=i` (**where**  $i > 1$ ):  $i$  goroutines compute hashes and update the map independently, using a mutex for thread safety.

- **Optional:** `-hash-workers=i -data-workers=j` (where  $i > j > 1$ ):  $i$  goroutines hash the BSTs, and  $j$  goroutines update the map, using either semaphores to control map access or multiple central managers with channels.

These flag combinations allow experimentation with different synchronization strategies, enabling performance analysis across parallel configurations while balancing concurrency, locking overhead, and map update efficiency.

### 3 Performance Metrics

The primary performance metrics in the BST equivalence program are:

#### 1. Hash Computation Time:

- Measures the time taken to compute and group hashes for all binary search trees (BSTs) in the dataset.
- Reflects the efficiency of the hashing phase and the impact of parallelism, such as the number of `hash-workers`.
- Recorded as `hashTime` in the program, this metric helps assess how effectively multiple threads manage the workload for hashing large datasets.

#### 2. Tree Comparison Time:

- Captures the time spent comparing trees with identical hashes to determine their equivalency.
- Measures the effectiveness of parallelization in the comparison phase, particularly relevant for larger datasets with numerous tree comparisons.
- Recorded as `compareTreeTime`, this metric evaluates the scalability of the comparison phase with various thread configurations.

#### 3. Hash Group Time: is a specific metric in the profiling of our binary search tree comparison program, defined and measured as follows:

- **Start Time:** This metric begins immediately after the completion of `hashTime`, marking the transition from hash computation to data handling.
- **Activities Included:** It encompasses the duration spent on processing the hash groups. This includes iterating over the hash map (`hash2bstId`) to either log, print, or otherwise handle the groups of tree indices grouped by their hash values.
- **End Time:** The measurement concludes once all activities related to hash group handling are completed, encapsulating the entire process of dealing with the organized data post-hashing.



## 4 Results and Analysis

### 4.1 Performance Data by Implementation Type with coarse.txt

#### Channel Implementation

Hash Workers	Data Workers	Hash Time	Compare Tree Time
1	1	0.02595379	0.00214175
2	1	0.01596462	0.00047217
4	1	0.01118975	0.00111958
6	1	0.01078404	0.00059246
8	1	0.00961188	0.00112513
10	1	0.00626321	0.00042579
12	1	0.00994050	0.00108975
14	1	0.01045579	0.00044925
16	1	0.01162029	0.00110000
18	1	0.01295263	0.00066967
20	1	0.01068587	0.00045408
24	1	0.00860700	0.00042721
28	1	0.01016667	0.00144492
32	1	0.01371433	0.00049987
36	1	0.00948421	0.00054458
48	1	0.01346096	0.00067604
56	1	0.01105333	0.00113542
64	1	0.01264725	0.00113696
100	1	0.01024737	0.00147663
128	1	0.01065667	0.00046650

## Mutex Implementation

Hash Workers	Data Workers	Hash Time	Compare Tree Time
1	1	0.02595379	0.00214175
2	2	0.01447592	0.00042912
4	4	0.01272117	0.00146558
6	6	0.01011950	0.00139383
8	8	0.01241625	0.00116621
10	10	0.00970850	0.00114625
12	12	0.00929833	0.00046092
14	14	0.00962158	0.00045333
16	16	0.00893458	0.00042150
18	18	0.00902125	0.00110467
20	20	0.00955975	0.00115712
24	24	0.01197079	0.00047704
28	28	0.01256971	0.00095183
32	32	0.01085713	0.00132167
36	36	0.01033062	0.00154800
48	48	0.01316883	0.00113954
56	56	0.00984496	0.00045733
64	64	0.00948858	0.00388983
100	100	0.01052967	0.00048442
128	128	0.01046562	0.00149058

## Semaphore Implementation

Hash Workers	Data Workers	Hash Time	Compare Tree Time
1	1	0.02595379	0.00214175
4	2	0.01068875	0.00047758
6	2	0.01029908	0.00126367
6	4	0.01070488	0.00282254
8	2	0.00894183	0.00046725
8	4	0.00906933	0.00043554
8	6	0.01333567	0.00069412
10	2	0.01194633	0.00115979
10	4	0.00984000	0.00118583
10	6	0.00999775	0.00047358
10	8	0.01247300	0.00113971
12	2	0.00955579	0.00111150
12	4	0.01165863	0.00057288
12	6	0.00905521	0.00046008
12	8	0.01204983	0.00101725
12	10	0.01086796	0.00114688
14	2	0.00963608	0.00107625
14	4	0.00942733	0.00112967
14	6	0.00960996	0.00107508
14	8	0.00953346	0.00047317
14	10	0.01016667	0.00118400
14	12	0.01356271	0.00044392
16	2	0.00954958	0.00120396
16	4	0.00901587	0.00115071
16	6	0.00992262	0.00108679
16	8	0.00987825	0.00107979
16	10	0.01128946	0.00108879
16	12	0.01047196	0.00108325
16	14	0.00949846	0.00139075
18	2	0.01190129	0.00110725
18	4	0.00928904	0.00117650
18	6	0.01373762	0.00168042
18	8	0.00975150	0.00112992
18	10	0.01015638	0.00124712
18	12	0.00987450	0.00109292
18	14	0.01014492	0.00108425
18	16	0.01007288	0.00107725
20	2	0.00955012	0.00112875
20	4	0.00932738	0.00120300
20	6	0.01026550	0.00045817
20	8	0.00913404	0.00046371
20	10	0.00932721	0.00108938
20	12	0.01037000	0.00122117
20	14	0.01326338	0.00130742
20	16	0.01043858 <sup>11</sup>	0.00045058

Hash Workers	Data Workers	Hash Time	Compare Tree Time
24	2	0.01151525	0.00119367
24	4	0.00936083	0.00044613
24	6	0.00874846	0.00045292
24	8	0.00996504	0.00137596
24	10	0.01023475	0.00045229
24	12	0.01001742	0.00108350
24	14	0.00972638	0.00116525
24	16	0.01100238	0.00108450
24	18	0.01152654	0.00110171
24	20	0.00956354	0.00139967
28	2	0.00931412	0.00143600
28	4	0.00960000	0.00047646
28	6	0.00979604	0.00107671
28	8	0.00979137	0.00107633
28	10	0.00963508	0.00045887
28	12	0.00912350	0.00108967
28	14	0.00947862	0.00046663
28	16	0.00910917	0.00042954
28	18	0.00926954	0.00045058
28	20	0.01307200	0.00126150
32	2	0.01346712	0.00137996
32	4	0.01279221	0.00054925
32	6	0.00979858	0.00045317
32	8	0.01021221	0.00108467
32	10	0.01037554	0.00200792
32	12	0.01034808	0.00128338
32	14	0.00996392	0.00046587
32	16	0.01154421	0.00046854
32	18	0.01159200	0.00118767
32	20	0.00929512	0.00045625
32	24	0.00969308	0.00063079
32	28	0.00969887	0.00140917
36	2	0.01473925	0.00113204
36	4	0.00935188	0.00044912
36	6	0.00955950	0.00046062
36	8	0.01004283	0.00115829
36	10	0.01283938	0.00108337
36	12	0.01037033	0.00174654
36	14	0.00983458	0.00118312
36	16	0.01198438	0.00046829
36	18	0.00906338	0.00045046
36	20	0.00885358	0.00043179
36	24	0.01049608	0.00057879
36	28	0.01169700	0.00126321
36	32	0.01008492	0.00061017

Hash Workers	Data Workers	Hash Time	Compare Tree Time
48	2	0.01018833	0.00046704
48	4	0.00996033	0.00045958
48	6	0.01049104	0.00109162
48	8	0.01026462	0.00129096
48	10	0.01183921	0.00152879
48	12	0.00908212	0.00108804
48	14	0.01001512	0.00302362
48	16	0.01161892	0.00119067
48	18	0.01028333	0.00138683
48	20	0.01034700	0.00045733
48	24	0.00988896	0.00110083
48	28	0.01046612	0.00123775
48	32	0.01004163	0.00347021
48	36	0.01031075	0.00046200
56	2	0.00922542	0.00045217
56	4	0.01010671	0.00121946
56	6	0.00989029	0.00128396
56	8	0.01028158	0.00128783
56	10	0.01058550	0.00122333
56	12	0.00981950	0.00109088
56	14	0.01036829	0.00131142
56	16	0.01231229	0.00101187
56	18	0.01167392	0.00243983
56	20	0.01082946	0.00136425
56	24	0.01092204	0.00123338
56	28	0.01044738	0.00111388
56	32	0.00992875	0.00118667
56	36	0.00976421	0.00046171
56	48	0.00965508	0.00134117
64	2	0.00887421	0.00042700
64	4	0.00961342	0.00143392
64	6	0.00996783	0.00108000
64	8	0.00967112	0.00115537
64	10	0.01040392	0.00145854
64	12	0.00855333	0.00051683
64	14	0.01024250	0.00108133
64	16	0.00961175	0.00047362
64	18	0.00973175	0.00137604
64	20	0.00883071	0.00042246
64	24	0.00965671	0.00109996
64	28	0.00943121	0.00126004
64	32	0.01054379	0.00129883
64	36	0.00965062	0.00045713
64	48	0.00995046	0.00111696
64	56	0.00978229	0.00129038

Hash Workers	Data Workers	Hash Time	Compare Tree Time
100	2	0.01436167	0.00072279
100	4	0.00977492	0.00124754
100	6	0.01198246	0.00125421
100	8	0.00988108	0.00120904
100	10	0.01246338	0.00105917
100	12	0.00913758	0.00043279
100	14	0.00996704	0.00109263
100	16	0.00978717	0.00112767
100	18	0.01000413	0.00330904
100	20	0.01069738	0.00132354
100	24	0.00991575	0.00045796
100	28	0.00973583	0.00128375
100	32	0.00979758	0.00046879
100	36	0.00987992	0.00047433
100	48	0.00975546	0.00119992
100	56	0.00980475	0.00109621
100	64	0.01075042	0.00046046
128	2	0.00942983	0.00045696
128	4	0.01011679	0.00079763
128	6	0.00983913	0.00046142
128	8	0.01006158	0.00177029
128	10	0.00949279	0.00124675
128	12	0.01002996	0.00131608
128	14	0.00951450	0.00046113
128	16	0.00984492	0.00047967
128	18	0.01042025	0.00269992
128	20	0.00972850	0.00046288
128	24	0.01015767	0.00114500
128	28	0.01037096	0.00170533
128	32	0.01061775	0.00162079
128	36	0.01008721	0.00130663
128	48	0.01080783	0.00127979
128	56	0.00980629	0.00108000
128	64	0.00914829	0.00045642
128	100	0.01004758	0.00113663

## 4.2 Performance Data by Implementation Type with fine.txt

### Channel Implementation with fine.txt

2	1	0.03030633	0.04595596
4	1	0.02999913	0.04238425
6	1	0.03613954	0.04128129
8	1	0.03403196	0.04561079
10	1	0.04717871	0.04610300
12	1	0.04039275	0.04157562
14	1	0.03805417	0.04133688
16	1	0.04032846	0.04166683
18	1	0.03180325	0.04168325
20	1	0.03132471	0.04166737
24	1	0.03989275	0.04151833
28	1	0.03895025	0.04143042
32	1	0.03703442	0.04277625
36	1	0.03710779	0.04331625
48	1	0.03673221	0.04253638
56	1	0.03791729	0.04285075
64	1	0.03509988	0.04289863
100	1	0.03724604	0.04246108
128	1	0.03550233	0.04210917

## Mutex Implementation with fine.txt

Hash Workers	Data Workers	Hash Time	Compare Tree Time
1	1	0.02595379	0.00214175
2	2	0.02841629	0.04293050
4	4	0.02407942	0.04274050
6	6	0.03490062	0.04607396
8	8	0.04324150	0.04399492
10	10	0.04625079	0.04157188
12	12	0.05325100	0.04148504
14	14	0.05477742	0.04143288
16	16	0.05364792	0.04160733
18	18	0.05538717	0.04125792
20	20	0.05138083	0.04357542
24	24	0.05327746	0.04150812
28	28	0.05635825	0.04200138
32	32	0.05693896	0.04611517
36	36	0.04278225	0.04385567
48	48	0.05482937	0.04320475
56	56	0.05680254	0.04220637
64	64	0.05609033	0.04253092
100	100	0.05711717	0.04347637
128	128	0.05104746	0.04279183



## Semaphore Implementation with fine.txt

Hash Workers	Data Workers	Hash Time	Compare Tree Time
4	2	0.03633617	0.04274392
6	2	0.04056079	0.04379625
6	4	0.04407029	0.07987642
8	2	0.04060600	0.04356071
8	4	0.05033217	0.04284621
8	6	0.05294988	0.04531692
10	2	0.05549325	0.05156400
10	4	0.05050579	0.04154479
10	6	0.05330175	0.04129350
10	8	0.05713929	0.04133775
12	2	0.03899954	0.04153275
12	4	0.04714100	0.04184879
12	6	0.08057408	0.04143346
12	8	0.06654762	0.04251925
12	10	0.05574237	0.04160212
14	2	0.03720508	0.04131725
14	4	0.04169067	0.04182388
14	6	0.06681583	0.04147162
14	8	0.08259996	0.04134112
14	10	0.06502679	0.04126275
14	12	0.05638979	0.04118258
16	2	0.03878392	0.04131733
16	4	0.05047588	0.04151196
16	6	0.06695387	0.04173213
16	8	0.08638667	0.04155408
16	10	0.07972821	0.04127338
16	12	0.06313208	0.04103892
16	14	0.05615542	0.04146412
18	2	0.03973279	0.04142371
18	4	0.04177575	0.04396896
18	6	0.05590450	0.04202162
18	8	0.09250392	0.04158392
18	10	0.07446546	0.04582171
18	12	0.06596554	0.04509908
18	14	0.05889912	0.04140404
18	16	0.05845842	0.04144929
20	2	0.03774408	0.04272363
20	4	0.04735771	0.04184796
20	6	0.04988225	0.04104167
20	8	0.07693900	0.04207433
20	10	0.09062183	0.04107175
20	12	0.08771663	0.04130158
20	14	0.06866400	0.04145546
20	16	0.06126800	0.04152400
20	18	0.05616283 <sup>17</sup>	0.04110279

## Semaphore Implementation with fine.txt

24	2	0.03963504	0.04160308
24	4	0.05049167	0.04140475
24	6	0.08333971	0.04621625
24	8	0.09379758	0.04444925
24	10	0.09588363	0.04145275
24	12	0.05679033	0.04163217
24	14	0.09155325	0.04140642
24	16	0.09007258	0.04182562
24	18	0.07753625	0.04151896
24	20	0.05726050	0.04142150
28	2	0.03951367	0.04309167
28	4	0.04318783	0.04156121
28	6	0.07610321	0.04136304
28	8	0.09474504	0.04127612
28	10	0.08116958	0.04126071
28	12	0.09039525	0.04207387
28	14	0.09702304	0.04127087
28	16	0.09308179	0.04229008
28	18	0.08889750	0.04131192
28	20	0.06976242	0.04118108
28	24	0.05803312	0.04141983
32	2	0.03909958	0.04137600
32	4	0.04304313	0.04120121
32	6	0.05122742	0.04124613
32	8	0.05173758	0.04101417
32	10	0.09491879	0.04132208
32	12	0.09661358	0.04132883
32	14	0.09844183	0.04143037
32	16	0.09110208	0.04117608
32	18	0.09558113	0.04441825
32	20	0.09745713	0.04570788
32	24	0.09209700	0.06299300
32	28	0.06130771	0.05901567
36	2	0.03954879	0.04355733
36	4	0.04079067	0.04379300
36	6	0.05136921	0.04255779
36	8	0.09722904	0.04225271
36	10	0.07550429	0.04461642
36	12	0.09794596	0.04255163
36	14	0.09645667	0.04276917
36	16	0.08732467	0.04250100
36	18	0.09646633	0.04350104
36	20	0.09484542	0.04368446
36	24	0.09490712	0.04292488
36	28	0.08359738	0.04344233
36	32	0.06036329	0.04357592

## Semaphore Implementation with fine.txt

48	2	0.03615037	0.04392258
48	4	0.04635933	0.04463487
48	6	0.07491154	0.04311138
48	8	0.08951412	0.04247779
48	10	0.09181804	0.04267221
48	12	0.09714225	0.04388729
48	14	0.09631375	0.04356758
48	16	0.08682404	0.04235092
48	18	0.09713183	0.04214933
48	20	0.09590692	0.04205842
48	24	0.09531954	0.04198575
48	28	0.09348463	0.04245954
48	32	0.07872525	0.04259617
48	36	0.05608400	0.04348462
56	2	0.03910979	0.04240179
56	4	0.04223654	0.04122912
56	6	0.06622058	0.04312158
56	8	0.06663967	0.04398500
56	10	0.09707262	0.04227492
56	12	0.07232679	0.04533042
56	14	0.09704358	0.04243292
56	16	0.08742812	0.04261133
56	18	0.09877579	0.04475746
56	20	0.09863508	0.04237292
56	24	0.09665633	0.04242812
56	28	0.09073338	0.04189321
56	32	0.08666842	0.04256167
56	36	0.08204925	0.04223321
56	48	0.06684363	0.05258717
64	2	0.06650750	0.04302479
64	4	0.04082717	0.04315112
64	6	0.05581946	0.04403471
64	8	0.07993121	0.04233387
64	10	0.09607138	0.04205562
64	12	0.09795954	0.04225763
64	14	0.09785933	0.04233996
64	16	0.09888637	0.04209492
64	18	0.09841883	0.04226875
64	20	0.08803017	0.04303121
64	24	0.09703088	0.04296025
64	28	0.06468858	0.04339575
64	32	0.08912267	0.04234929
64	36	0.08054063	0.04280213
64	48	0.04375242	0.04312513
64	56	0.05644329	0.04495763

## Semaphore Implementation with fine.txt

100	2	0.03892654	0.04328808
100	4	0.04613908	0.04309012
100	6	0.07656821	0.04276783
100	8	0.09460117	0.04307250
100	10	0.09550996	0.04222904
100	12	0.09765346	0.04241117
100	14	0.09951671	0.04318579
100	16	0.08687017	0.04247354
100	18	0.09943975	0.04217292
100	20	0.09601850	0.04417833
100	24	0.09361387	0.04229479
100	28	0.09368654	0.04214788
100	32	0.08692683	0.04507308
100	36	0.08833396	0.04217908
100	48	0.06445162	0.04612767
100	56	0.07617638	0.04381258
100	64	0.06194413	0.04167967
128	2	0.03852796	0.04462762
128	4	0.04674971	0.04285463
128	6	0.06866633	0.05114233
128	8	0.09802329	0.04365358
128	10	0.09472650	0.04231717
128	12	0.08419675	0.04228262
128	14	0.09669779	0.04277171
128	16	0.09328825	0.04226246
128	18	0.09116354	0.04222671
128	20	0.09345721	0.04383267
128	24	0.09676517	0.04209258
128	28	0.08142325	0.04282725
128	32	0.07719925	0.04246213
128	36	0.08276821	0.04381946
128	48	0.07333517	0.04204200
128	56	0.08884438	0.04280550
128	64	0.06414283	0.04253804
128	100	0.05932346	0.04231342

## 4.3 Speedup Graphs

### 4.3.1 Hash Computation Time

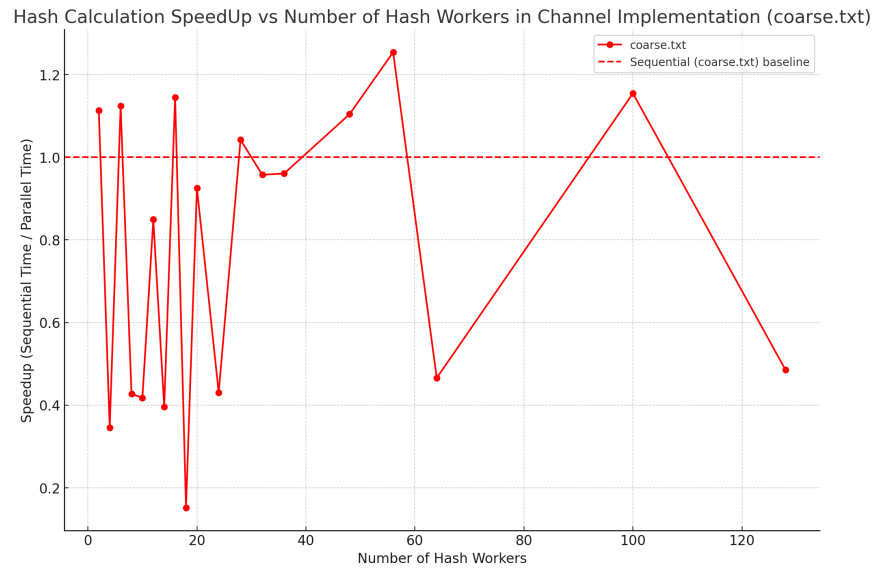


Figure 1: Hashworkers vs SpeedUp using channel for coarse.txt




Figure 2: Hashworkers vs SpeedUp for coarse.txt using mutex for coarse.txt

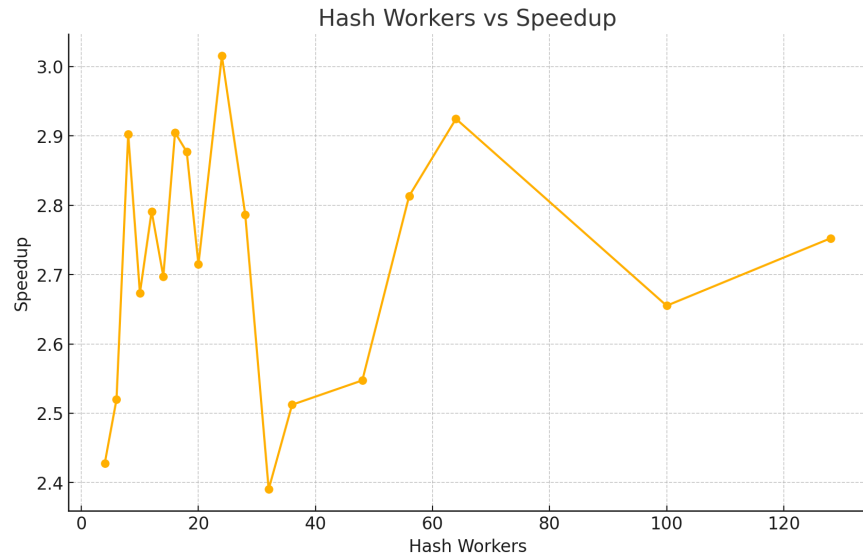


Figure 3: Hashworkers vs SpeedUp using semaphore for coarse.txt

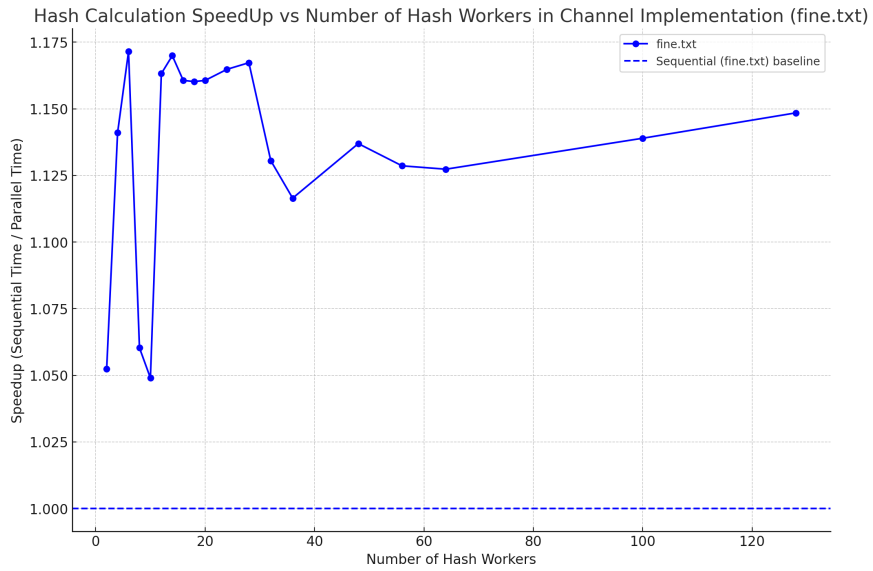


Figure 4: Hashworkers vs SpeedUp using channel for fine.txt

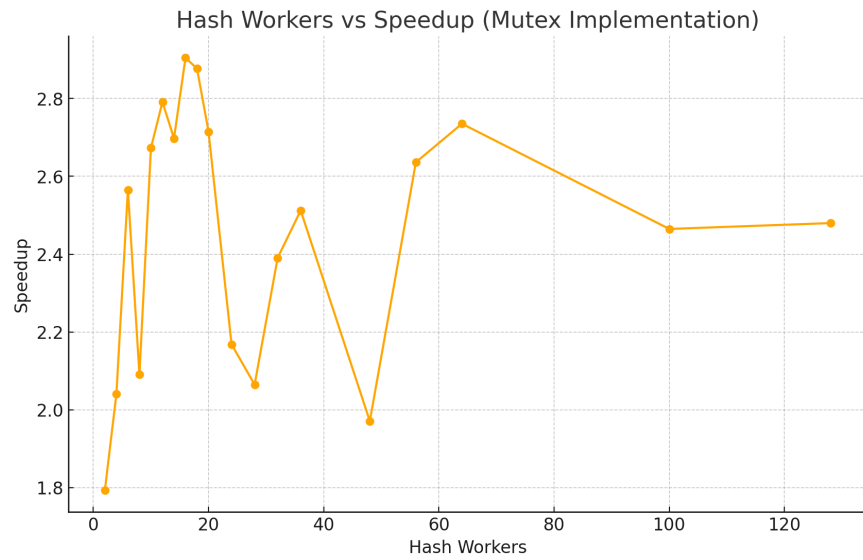


Figure 5: Hashworkers vs SpeedUp using mutex for fine.txt

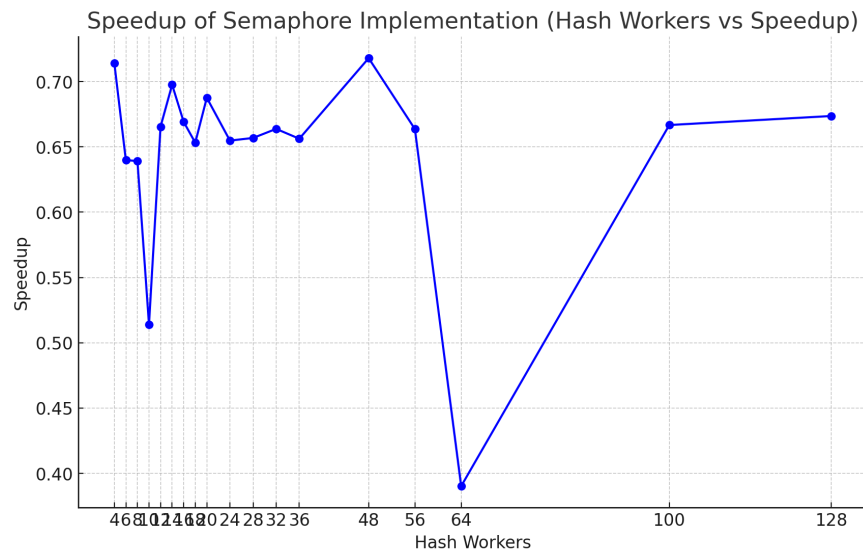


Figure 6: Hashworkers vs SpeedUp using semaphore for fine.txt

### 4.3.2 Tree Comparison Time

comp-workers	compareTreeTime
1	0.00057900
2	0.00033346
4	0.00026146
6	0.00024288
7	0.00042671
8	0.00278529
9	0.00022721
10	0.00054133
12	0.00021196
14	0.00043604
16	0.00024308
18	0.00136396
20	0.00059437
24	0.00138379
28	0.00024963
32	0.00022996
36	0.00077021
48	0.00226400
56	0.00025721
64	0.00027312
100	0.00035492
128	0.00075346

Table 1: Comparison of Tree Times by Number of Workers for coarse.txt



<b>comp-workers</b>	<b>compareTreeTime</b>
1	0.04143842
2	0.05584687
4	0.04198763
6	0.03719104
7	0.03862433
8	0.03703254
9	0.03675567
10	0.03453496
12	0.04155637
14	0.03429696
16	0.04181454
18	0.03734808
20	0.04253371
24	0.03834558
28	0.03861767
32	0.04101271
36	0.04212392
48	0.05031371
56	0.04995446
64	0.05237017
100	0.06949183
128	0.07367017

Table 2: Comparison Workers and Tree Comparison Time for fine.txt

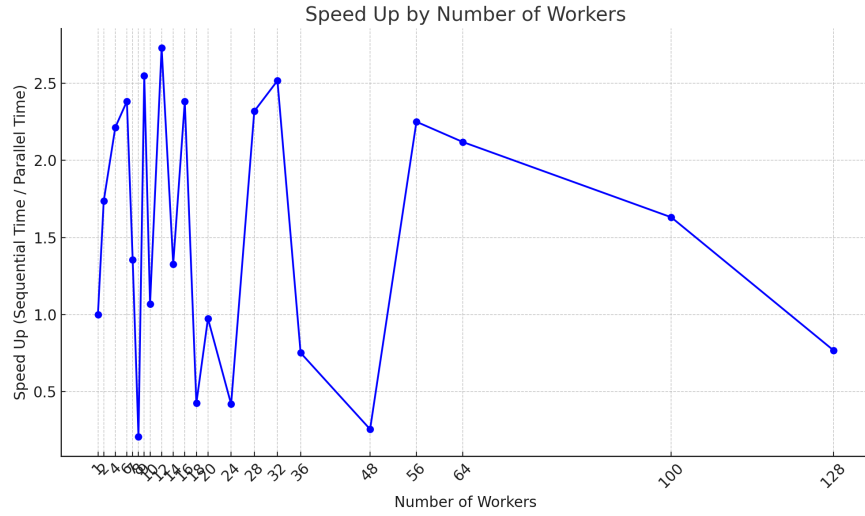


Figure 7: comp workers vs SpeedUp for coarse.txt

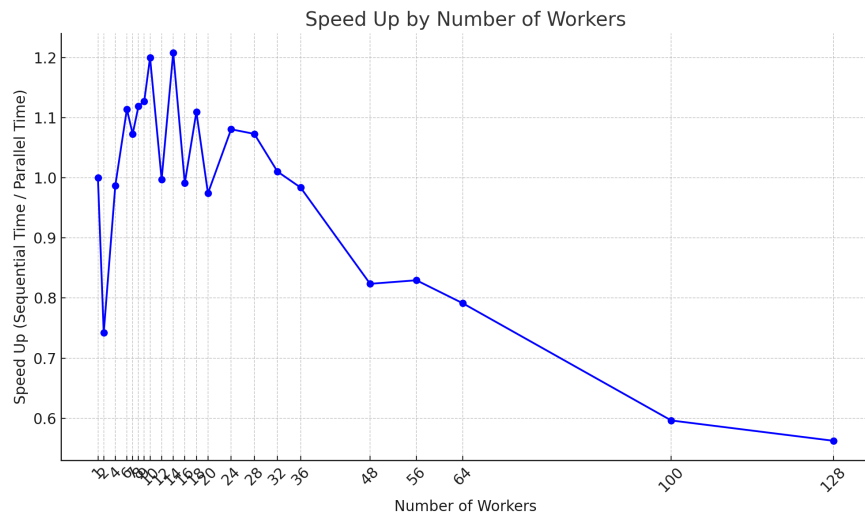


Figure 8: comp workers vs SpeedUp for fine.txt

## 4.4 Correctness Verification

The output is same with expected output provided by one of the student.

```

~/Desktop/Parallel systems/projects/lab3 (master*) » ./myprogram -input=./input/coarse.txt -comp-workers=7
hashTime: 0.02401667
31: 0 3 9 13 15 16 17 22 24 27 28 30 32 35 37 43 45 46 48 49 52 54 60 64 65 66 74 78 79 86 88 89 90 91 94 98
191: 1 2 4 5 6 7 8 10 11 12 14 18 19 20 21 23 25 26 29 31 33 34 36 38 39 40 41 42 44 47 50 51 53 55 56 57 58 59
61 62 63 67 68 69 70 71 72 73 75 76 77 80 81 82 83 84 85 87 92 93 95 96 97 99
hashGroupTime: 0.00003017
compareTreeTime: 0.00056929
group 0: 10 76
group 1: 17 65
group 2: 19 21
group 3: 30 88
group 4: 41 83
group 5: 56 95
group 6: 4 71
group 7: 25 36
group 8: 26 34 63
group 9: 40 62
group 10: 81 85 97
group 11: 14 50 57
group 12: 91 98
group 13: 20 53 75
group 14: 16 27 35 78
group 15: 8 80
group 16: 59 77
group 17: 67 72
group 18: 69 70
group 19: 0 9 45 86
group 20: 5 31
group 21: 11 18 29 96
group 22: 38 93 99
group 23: 6 47
group 24: 23 33 51
group 25: 42 68 92
group 26: 74 79
group 27: 3 60
group 28: 12 73 87
group 29: 28 54 94
group 30: 1 55
group 31: 37 52 89
group 32: 15 24 49

```

Figure 9: correctness validation

## 5 Discussion

### 5.1 Channel Implementation

- **Concurrent Processing:** As we can see from the graphs, Channel-based implementation excels with larger input size, the workload is large enough to benefit from multiple workers processing in parallel, significantly reducing the total computation time, and result is more consistent.
- **Reduced Contention:** This approach minimizes direct contention for shared resources by using channels for communication, which can lead to improved throughput on larger datasets.
- **Overhead and Blocking:** The overhead associated with managing channels and potential blocking on channel operations can become a bottleneck, especially with smaller or unevenly distributed workloads.

### 5.2 Mutex Implementation

- **Direct Map Access:** Mutex implementation allows workers to directly update the shared map, potentially reducing the complexity of data handling compared to channel-based approaches.
- **Contention Issues:** The use of a single mutex can lead to significant contention, especially with a high number of concurrent workers, which can negate the benefits of parallelism. That is why with higher number of workers, the speedup drops.
- **Scalability Concerns:** While simpler to implement, the scalability of this method is limited by the mutex's ability to handle high contention levels effectively. Which is why we can see from the graphs, speedup for mutex impl for coarse.txt and fine.txt does not differ too much.

### 5.3 Semaphore Implementation

- **Fine-grained Locking:** By using semaphores or fine-grained locking strategies, this implementation can offer better performance scaling by reducing the contention seen in single-mutex setups.
- **Complexity and Overhead:** However, managing multiple locks or semaphores introduces additional complexity and synchronization overhead, which might not always translate to proportionate performance gains. Which is why we can see for larger input of semaphore speedup graph, there isn't much speedup comparing to sequential implementation. But for medium sized input file, we do observe some kind of speedup.
- **Optimal for Medium to Large Workloads:** Semaphore strategies are most effective for medium to large workloads where the overhead can be justified by significant performance improvements over simpler locking mechanisms.

## 5.4 Comparison Workers and Speedup Analysis

We do not observe liner speed up, adding more comp workers suffers from some overheads introduced in task scheduling, worker synchronization. While it does speedup better for coarse.txt with medium sized input, it does not scale well for the larger size input.

Tasks (tree comparisons) are distributed to compWorkers via channels. This involves dynamically allocating pairs of trees to be compared and managing these tasks through channels. Each time a task is sent through a channel, there's an overhead associated with context switching and managing these communications. As the number of workers increases, the overhead associated with scheduling and distributing these tasks also grows.

Depending on how the tree comparisons are batched and distributed, some workers might end up with more work than others. This imbalance can cause some workers to be idle while others are still processing, reducing overall efficiency.

### 5.4.1 Impact of Number of Workers

- **Few Workers:** With too few workers, the potential for parallel processing is not fully realized, often leading to underutilization of available processing resources.
- **Many Workers:** Conversely, too many workers can lead to excessive context switching, increased contention for shared resources, and overall inefficiency, particularly when the number of workers exceeds the number of processing cores.
- **Optimal Number of Workers:** The optimal number of workers typically lies between these two extremes and is dependent on the specific characteristics of the hardware and the nature of the tasks being processed.

### 5.4.2 Recommendations for Optimization

- **Dynamic Worker Allocation:** Implementing a dynamic worker allocation strategy that adjusts the number of workers based on the workload can help in achieving near-optimal performance across different scenarios.
- **Workload Balancing:** Efforts should be made to ensure that the workload is evenly distributed among the available workers to avoid bottlenecks and idle resources.
- **Hardware Considerations:** Understanding the underlying hardware capabilities, such as the number of cores and the memory hierarchy, can guide the configuration of the parallel execution environment to better align with the systems strengths.

This analysis underscores the importance of carefully selecting the number of comparison workers to maximize the efficiency of parallel tree comparisons. The insights derived from speedup metrics assist in fine-tuning the concurrency model to better harness the computational power available, ensuring that the parallel algorithm performs optimally under varying operational conditions.

## 5.5 General Observations

- **Performance with Larger Inputs:**

- parallel implementations tend to perform better with larger input sizes where the benefits of concurrent execution outweigh the overhead of managing parallel processes, even though this observation may not hold true for all the implementations.

- **Performance with Smaller Inputs:**

- Challenges such as setup cost, task granularity, and sensitivity to system conditions often result in less consistent performance with smaller datasets.

- **Optimization Strategies:**

- Dynamic adjustment of the concurrency level and employing hybrid strategies that switch between sequential and parallel processing based on the input size could optimize performance across varying workloads.

## 6 Conclusion and Improvement Ideas

The analysis of BST comparison reveals significant insights into the behavior of different concurrency models under varying workloads. While the channel-based implementation showcases its strength in environments with larger datasets by efficiently minimizing contention and maximizing throughput, it struggles with overheads in smaller workloads. Conversely, mutex and semaphore implementations highlight trade-offs between simplicity and scalability, with mutexes facing bottlenecks in high contention scenarios and semaphores introducing complexity that may not always yield proportional benefits. The findings underscore the necessity of a balanced approach in worker allocation and task distribution to harness full computational capabilities. Optimizing the number of workers and refining synchronization mechanisms based on empirical data and hardware capabilities can significantly enhance the performance and scalability of the BST comparison application across diverse operational conditions.

## References

### Technical Sources:

- GoLang document
- Concurrency Synchronization Techniques Provided in the sync Standard Package
- gos-extended-concurrency-semaphores
- understanding-and-implementing-the-semaphore-pattern-in-go